



FastUp: Fast TCAM Update for SDN Switches in Datacenter Networks

Ying Wan¹, Haoyu Song², Hao Che³, Yang Xu⁴, Yi Wang⁵, Chuwen Zhang¹, Zhijun Wang³, Tian Pan⁸, Hao Li⁷, Hong Jiang³, Chengchen Hu^{6,7}, Bin Liu¹

¹Tsinghua University, China, ²Futurewei Technologies, USA

³University of Texas at Arlington, USA, ⁴Fudan University, China

⁵Southern University of Science and Technology, China,

⁶Xilinx, Singapore, ⁷Xi'an Jiaotong University, China

⁸Beijing University of Posts and Telecommunications, China

Email: wany16@mails.Tsinghua.edu.cn



清华大学
Tsinghua University

Outline

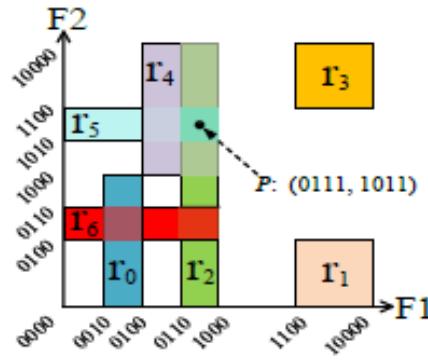
- Background and related work
- The principle of *FastUp*:
- Analysis of the optimal TCAM update
- Performance evaluation
- Conclusion

Background —— TCAM for rule tables

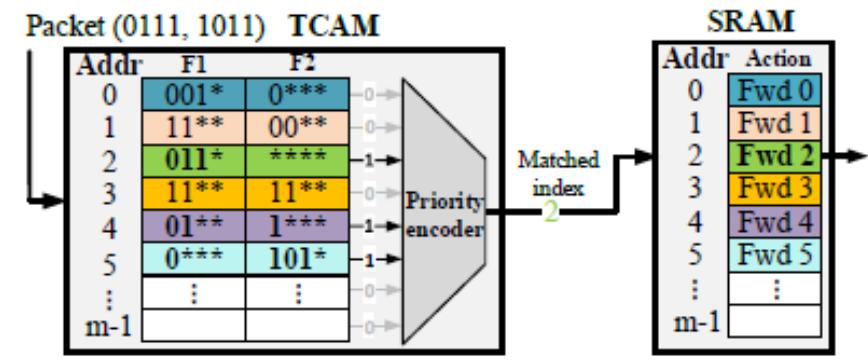
- The *de facto* industry standard for rule tables
 - line-speed lookup speed: compare a key with all rules in parallel
 - flexible matching pattern: LPM, EM, and RM
- The placed rules are arranged in priority order
 - TCAM always returns the first matched rule (i.e., with the lowest physical address)
 - It is the highest-priority matched rule that counts
 - E.g., packet p matches r2, r4, and r5, but the lowest one r2 is returned
- A new rule insertion incurs the moves of existing rules
 - R6 overlaps with r0 and r2. r6 has a higher priority than r2 and has a lower priority than r0
 - Up to 10 ms on P4 switch to insert a new rule when the TCAM capacity is set to 4K entries

Rule	Pri	Match		Action
		F1	F2	
r ₀	9	001*	0***	Fwd 0
r ₁	7	11**	00**	Fwd 1
r ₂	6	011*	****	Fwd 2
r ₃	5	11**	11**	Fwd 3
r ₄	2	01**	1***	Fwd 4
r ₅	1	0***	101*	Fwd 5

(a) Flow table.



(b) Rule overlapping.



(c) Implementation of the flow table in TCAM and SRAM.

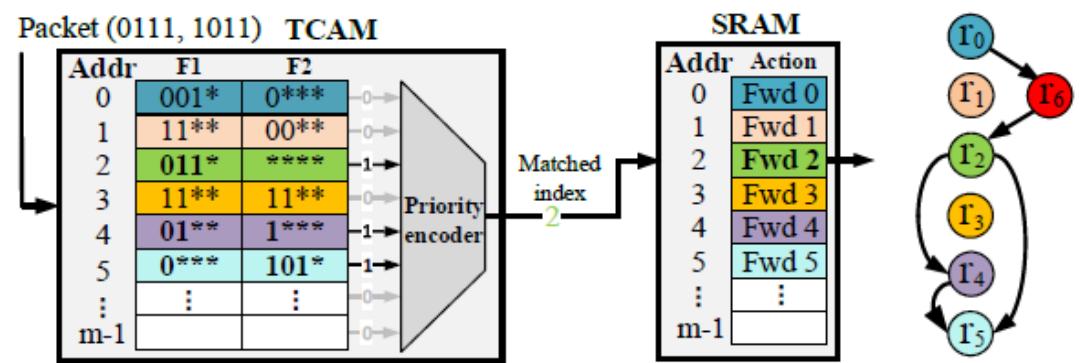
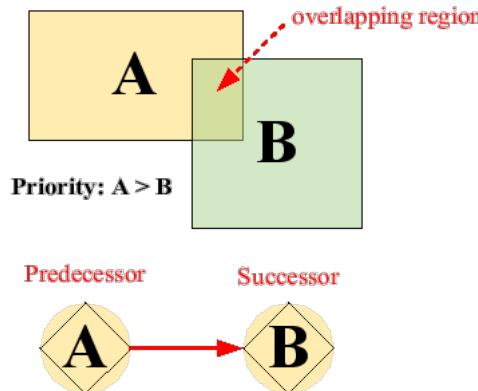
Background —— TCAM update problem

- **Rule relation graph**

- each node represents a rule and each edge represents the overlapping relationship
- each rule must be placed above all its descendants and below all its ascendants.
 - R's predecessor: the descendant of r that is the closest to r
 - R's successor: the descendant of r that is the closest to r

- **Rule moving strategy**

- A rule can be inserted between its predecessor and its successor
- A rule r can be moved downward directly to its successor
- A rule r can be moved downward to any entry between it and its successor.



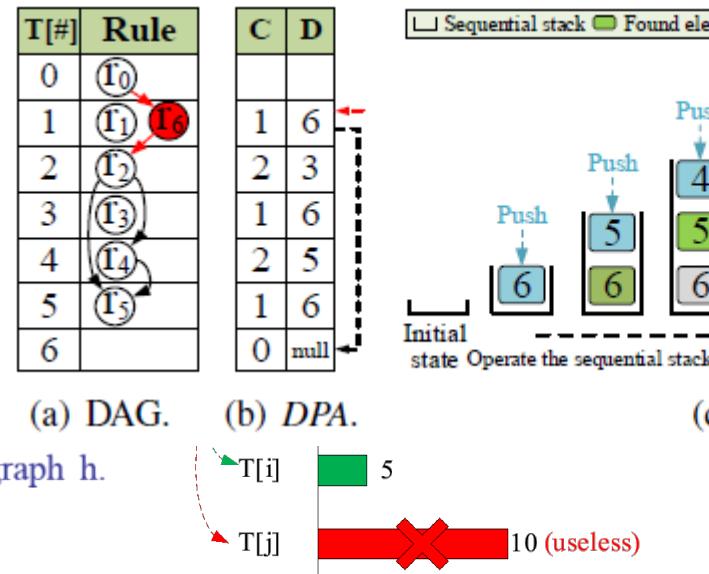
SSA——sequential-stack based algorithm

- Definition of moving cost
 - $C[i]$: the minimum rule moves to relocate the rule in the i -th TCAM entry $T[i]$
 - $\text{Succ}(i).\text{addr}$: the address of the entry of the successor of the rule in $T[i]$
 - $C[i] = \min_{j \in (i, \text{succ}(i).\text{addr})} \{C[j]\} + 1$
 - $D[i]$ records the destination entry to relocate the rule currently placed in $T[i]$
 - Calculation of $C[:]/D[:]$ is typical dynamic programming (DP) process.

- Optimization opportunity

- If $j > i$ and $C[j] > C[i]$, $T[j]$ will never be the best candidate
 - Dynamic identify and remove the useless entries
 - Only Find the best candidate among the entries potentially candidates
- Sequential stack S

- Time complexity: $O(m^2) \rightarrow O(m * lgh)$
 - The elements of S is in strict decreasing order.
 - If $T[i]$ is placed in $S[p]$, then $C[i] = p$.
 - The length of S will never exceed the diameter of rule graph h.
- Memory footprint: $O(m) \rightarrow O(h)$
 - If $T[i]$ is placed in $S[p]$, then $D[i] = S[p-1]$.
 - The best candidate of the update rule must stay in S after calculation.



RCA—Rule chain based algorithm for reorder resolution

- Definition of reorder problem
 - SSA assumes that a new rule ru can always find one or more candidate locations.
 - If $\text{succ}(\text{ru}).\text{addr} < \text{pred}(\text{ru}).\text{addr}$, the reorder problem happens.
 - E.g., $r3 \rightarrow r6 \rightarrow r0$
 - Relocate the out-of-order rules until $\text{succ}(\text{ru}).\text{addr} > \text{pred}(\text{ru}).\text{addr}$
- Reorder resolution
 - Why not choose SSA:
 - Time complexity and infinite loop
 - RCA:
 - Move one of out-of-order two rules and keep the other one in place
 - $r0 \rightarrow r0.\text{succ} \rightarrow r0.\text{succ}.\text{succ}$
 - Guarantee the reorder resolution
 - in one round, if the reorder is not resolved, the gap between the rules is reduced.

T[#]	Rule
0	T_0
1	T_1
2	T_2
3	T_3
4	T_4
5	T_5
6	

(c) RCA.

T[#]	Rule	C	D
0	T_0	2	3
1	T_1	3	2
2	T_2	2	3
3	T_3	1	6
4	T_4	2	5
5	T_5	1	6
6		0	null

(a) DPA.

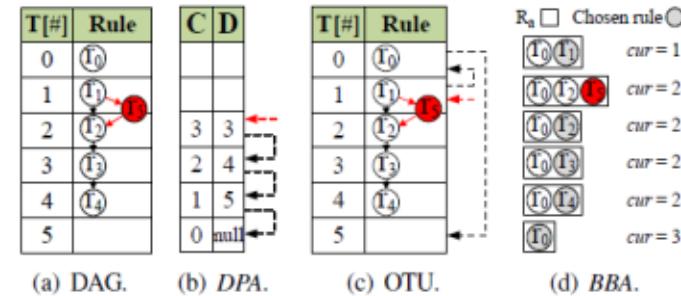
T[#]	Rule	C	D
0		0	null
1	T_1	1	0
2	T_2	2	1
3	T_0	1	0
4	T_4	2	3
5	T_5	3	4
6	T_3	2	3

(b) DPA.

BBA—branch-and-bound algorithm for optimal TCAM update

- Definition of optimal TCAM update (OTU)
 - The solution with the theoretically smallest rule moves
- Why exist algorithms fail to find the optimal solution
 - The rules are moved in the fixed and same direction
- The difficulty in achieving OTU
 - any topological order of the DAG is a feasible TCAM placement and vice versa.
 - N_r, m, N_{TT} , and N_L are the number of rules, TCAM entries, topological orders and feasible TCAM layouts

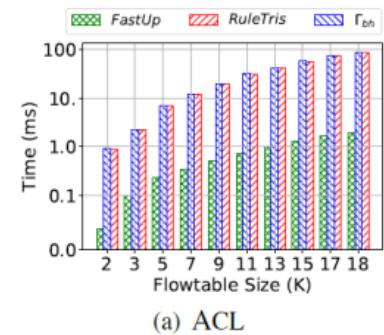
$$N_L = N_\pi * \binom{n}{m} = N_\pi * \frac{m!}{n!(m-n)!}$$



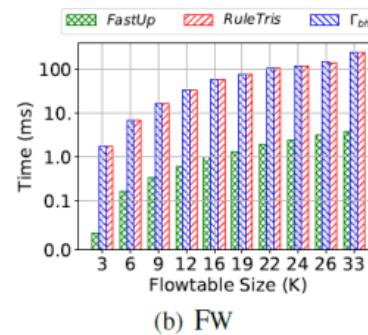
- just finding the number of all topological orders N_{TT} has been proven to be NP-hard.
- Why need to find the OTU
 - The degree of optimality, j , for a Design Under Test (DUT) is defined as $\lambda_{\text{DUT}} = \frac{N_{\text{OTU}}}{N_{\text{DUT}}} \times 100\%$
 - can be used to guide further algorithm optimizations.
- How to find the OTU
 - BBA processes each entry from top to bottom
 - BBA tries to place each available rule in the current entry, or leave it empty.
 - BBA avoids searching the space when the cumulative update cost exceeds the found OTU

Performance evaluation

- Experiment setup
 - Compare objects: RuleTris, Γ_{bh} , Γ_{down}
 - Testbed: a programmable OpenFlow switch——ONetSwitch
 - Dataset : Access Control List (ACL) and Firewall (FW), generated by ClassBench [59].
- Experimental results :
 - Metric: Compute time, interrupt time, reorder efficiency, and optimality degree



(a) ACL



(b) FW

$S_2(k)$	#Updates	Average time (ms)			Maximal time (ms)		
		FastUp	Γ_{bh}	Γ_{down}	FastUp	Γ_{bh}	Γ_{down}
3.4	340	0.62	0.65	2.31	1.2	1.8	6.6
7.4	740	0.64	0.78	3.16	2.4	3.6	10.8
11.4	1140	0.65	0.85	4.23	2.4	4.2	10.8
15.4	1540	0.70	0.86	4.57	3	4.8	11.4
18.6	1860	0.72	0.95	4.78	3.6	6	18

Case	Probability	Interrupt time (ms)					
		FastUp		BBA		λ_{FastUp}	
		Avg	Max	Avg	Max	Avg	Max
Normal	97.23%	1.29	4.2	1.24	3.0	96%	71%
Reorder	2.77%	5.20	12.6	2.30	3.0	44%	23%
Mixed	100%	1.40	12.6	1.27	3.0	90%	23%

Conclusion

- FastUp optimizes both the compute time and interrupt time
- FastUp solves the reorder problem efficiently
- FastUp is close to the optimal TCAM update

Thank You!

Q & A